

Tensforwflow example with two hidden layers. The MNIST dataset

Out[1]: [Click here to show or hide your raw code.](#)

Q1: (hidden layer size)



Model: Try a hidden layer size of 50

HyperParamters:
BUFFER SIZE =10000, BATCH SIZE = 100, NUM_EPOCHS = 5
Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 50)	39250
dense_1 (Dense)	(None, 50)	2550
dense_2 (Dense)	(None, 10)	510

Total params: 42,310
Trainable params: 42,310
Non-trainable params: 0

None



Train Accuracy = 0.972
Train Loss = 0.095

Validation Accuracy = 0.972
Validation Loss = 0.102

Test Accuracy = 0.970
Test Loss = 0.104

Taining Time = 7.903 second



Model: Try a hidden layer size of 200

HyperParameters:
BUFFER SIZE =10000, BATCH SIZE = 100, NUM_EPOCHS = 5
Model: "sequential_1"

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 784)	0
dense_3 (Dense)	(None, 200)	157000
dense_4 (Dense)	(None, 200)	40200
dense_5 (Dense)	(None, 10)	2010

Total params: 199,210
Trainable params: 199,210
Non-trainable params: 0

None



Train Accuracy = 0.988
Train Loss = 0.041

Validation Accuracy = 0.987
Validation Loss = 0.045

Test Accuracy = 0.976
Test Loss = 0.080

Taining Time = 7.000 second



Model: Try a Custom hidden layer size

HyperParameters:
BUFFER SIZE =10000, BATCH SIZE = 100, NUM_EPOCHS = 5
Model: "sequential_2"

Layer (type)	Output Shape	Param #
flatten_2 (Flatten)	(None, 784)	0
dense_6 (Dense)	(None, 1000)	785000
dense_7 (Dense)	(None, 1000)	1001000
dense_8 (Dense)	(None, 10)	10010

Total params: 1,796,010
Trainable params: 1,796,010
Non-trainable params: 0

None



Train Accuracy = 0.991
Train Loss = 0.029

Validation Accuracy = 0.989
Validation Loss = 0.040

Test Accuracy = 0.979
Test Loss = 0.073

Taining Time = 7.385 second

Answering Q1: (hidden layer size)

1. Try a hidden layer size of 200. How does the validation accuracy of the model change? What about the time it took the algorithm to train? Can you find a hidden layer size that does better?

How does the validation accuracy of the model change?

When increase the hidden layer from 50 to 200, the validation accuracy also increased from 0.970 to 0.988

What about the time it took the algorithm to train?

Train time increased from 7.396 to 7.856 seconds ¶

Can you find a hidden layer size that does better?

hidden layer size = 1000, I can notice that increasing the hidden layer is increaseing the validation accuracy and testing accuracy, but if you keep increasing hidden layer more than 1000 the accuracy almost the same but the training time increaseing, until you the model will overfit and the accuracy will start in increaseing.

Q2: (The Depth of hidden layers)



Model: Add another hidden layer

HyperParameters:
BUFFER SIZE = 10000, BATCH SIZE = 100, NUM_EPOCHS = 5
Model: "sequential_3"

Layer (type)	Output Shape	Param #
flatten_3 (Flatten)	(None, 784)	0
dense_9 (Dense)	(None, 1000)	785000
dense_10 (Dense)	(None, 1000)	1001000
dense_11 (Dense)	(None, 1000)	1001000
dense_12 (Dense)	(None, 10)	10010

Total params: 2,797,010
Trainable params: 2,797,010
Non-trainable params: 0

None



Train Accuracy = 0.988
Train Loss = 0.039

Validation Accuracy = 0.987
Validation Loss = 0.044

Test Accuracy = 0.980
Test Loss = 0.080

Taining Time = 7.768 second

Answering Q2: (The Depth of hidden layers)

2. The *depth* of the algorithm. Add another hidden layer to the algorithm. This is an extremely important exercise! How does the validation accuracy change? What about the time it took the algorithm to train?

How does the validation accuracy change?

When Adding another hidden layer, the validation accuracy decreased from 0.987 to 0.982

What about the time it took the algorithm to train?

Train time increased from 7.678 to 8.179 seconds

Q3: (Width and Depth of the Model)



Model: (5 Hidden layers, 1000 width)

HyperParameters:
BUFFER SIZE =10000, BATCH SIZE = 100, NUM_EPOCHS = 5
Model: "sequential_4"

Layer (type)	Output Shape	Param #
flatten_4 (Flatten)	(None, 784)	0
dense_13 (Dense)	(None, 1000)	785000
dense_14 (Dense)	(None, 1000)	1001000
dense_15 (Dense)	(None, 1000)	1001000
dense_16 (Dense)	(None, 1000)	1001000
dense_17 (Dense)	(None, 1000)	1001000
dense_18 (Dense)	(None, 10)	10010

Total params: 4,799,010
Trainable params: 4,799,010
Non-trainable params: 0

None



Train Accuracy = 0.985
Train Loss = 0.053

Validation Accuracy = 0.975
Validation Loss = 0.100

Test Accuracy = 0.964
Test Loss = 0.136

Taining Time = 8.959 second

3. The *width and depth* of the algorithm. Add as many additional layers as you need to reach 5 hidden layers. Moreover, adjust the width of the algorithm as you find suitable. How does the validation accuracy change? What about the time it took the algorithm to train?

How does the validation accuracy of the model change?

When use (5 hidden layers, 1000 width) the validation accuracy also increased from 0.982 to 0.987.

What about the time it took the algorithm to train?

Train time increased from 8.179 to 9.193 seconds

Q4: (Fiddle with the activation functions.)



Model: (use Sigmoid Activation Funcation)

HyperParameters:
BUFFER SIZE =10000, BATCH SIZE = 100, NUM_EPOCHS = 5
Model: "sequential_5"

Layer (type)	Output Shape	Param #
flatten_5 (Flatten)	(None, 784)	0
dense_19 (Dense)	(None, 1000)	785000
dense_20 (Dense)	(None, 1000)	1001000
dense_21 (Dense)	(None, 1000)	1001000
dense_22 (Dense)	(None, 1000)	1001000
dense_23 (Dense)	(None, 1000)	1001000
dense_24 (Dense)	(None, 10)	10010

Total params: 4,799,010
Trainable params: 4,799,010
Non-trainable params: 0

None



Train Accuracy = 0.985
Train Loss = 0.051

Validation Accuracy = 0.979
Validation Loss = 0.075

Test Accuracy = 0.973
Test Loss = 0.097

Taining Time = 8.934 second

4. Fiddle with the activation functions. Try applying sigmoid transformation to both layers. The sigmoid activation is given by the string 'sigmoid'.

How does the validation accuracy of the model change?

use Sigmoid in 2 hidden layers, the validation accuracy also decreased from 0.987 to 0.986. Test Accuracy decreased from 0.979 to 0.977

What about the time it took the algorithm to train?

Train time decreased from 9.193 to 9.190 seconds

Q5: (Fiddle with the activation functions.)

Model: (use Tanh Activation Funcation)

HyperParameters:

BUFFER SIZE =10000, BATCH SIZE = 100, NUM_EPOCHS = 5

Model: "sequential_6"



Layer (type)	Output Shape	Param #
flatten_6 (Flatten)	(None, 784)	0
dense_25 (Dense)	(None, 1000)	785000
dense_26 (Dense)	(None, 1000)	1001000
dense_27 (Dense)	(None, 1000)	1001000
dense_28 (Dense)	(None, 1000)	1001000
dense_29 (Dense)	(None, 1000)	1001000
dense_30 (Dense)	(None, 10)	10010

Total params: 4,799,010

Trainable params: 4,799,010

Non-trainable params: 0

None



Train Accuracy = 0.985

Train Loss = 0.055

Validation Accuracy = 0.982

Validation Loss = 0.068

Test Accuracy = 0.975

Test Loss = 0.103

Taining Time = 8.911 second

5. Fiddle with the activation functions. Try applying sigmoid transformation to both layers. The sigmoid activation is given by the string 'sigmoid'.

How does the validation accuracy of the model change?

use Sigmoid in 2 hidden layers, the validation accuracy also decreased from 0.986 to 0.980. Test Accuracy decreased from 0.977 to 0.972

What about the time it took the algorithm to train?

Train time increased from 9.190 to 9.380 seconds

Q6: (Adjust the batch size)



Model: (use batch size 10000)

HyperParameters:
BUFFER SIZE =10000, BATCH SIZE = 10000, NUM_EPOCHS = 5
Model: "sequential_7"

Layer (type)	Output Shape	Param #
flatten_7 (Flatten)	(None, 784)	0
dense_31 (Dense)	(None, 1000)	785000
dense_32 (Dense)	(None, 1000)	1001000
dense_33 (Dense)	(None, 1000)	1001000
dense_34 (Dense)	(None, 1000)	1001000
dense_35 (Dense)	(None, 1000)	1001000
dense_36 (Dense)	(None, 10)	10010

Total params: 4,799,010
Trainable params: 4,799,010
Non-trainable params: 0

None



Train Accuracy = 0.933
Train Loss = 0.226

Validation Accuracy = 0.945
Validation Loss = 0.194

Test Accuracy = 0.940
Test Loss = 0.189

Taining Time = 6.075 second

6. Adjust the batch size. Try a batch size of 10000. How does the required time change? What about the accuracy?

How does the required time change?

Train time decreased from 9.193 to 6.044 seconds

What about the accuracy?

validation Accuracy decreased from 0.987 to 0.948, Test Accuracy decreased from 0.979 to 0.945,

Q7: (Adjust the batch size)



Model: (use batch size 1)

HyperParameters:
BUFFER SIZE =10000, BATCH SIZE = 1, NUM_EPOCHS = 5
Model: "sequential_8"

Layer (type)	Output Shape	Param #
flatten_8 (Flatten)	(None, 784)	0
dense_37 (Dense)	(None, 1000)	785000
dense_38 (Dense)	(None, 1000)	1001000
dense_39 (Dense)	(None, 1000)	1001000
dense_40 (Dense)	(None, 1000)	1001000
dense_41 (Dense)	(None, 1000)	1001000
dense_42 (Dense)	(None, 10)	10010

Total params: 4,799,010
Trainable params: 4,799,010
Non-trainable params: 0

None



Train Accuracy = 0.957
Train Loss = 0.225

Validation Accuracy = 0.975
Validation Loss = 0.155

Test Accuracy = 0.969
Test Loss = 0.201

Taining Time = 449.658 second

7. Adjust the batch size. Try a batch size of 1. How does the required time change? What about the accuracy?

How does the required time change?

Train time significantly Increased from 9.193 to 449.6 seconds

What about the accuracy?

validation Accuracy significantly decreased from 0.987 to 0.975 . , Test Accuracy decreased from 0.979 to 0.969.

Q8: (Adjust the Learning rate)



Model: (use Learning rate 0.0001)

HyperParameters:
BUFFER SIZE =10000, BATCH SIZE = 1000, NUM_EPOCHS = 5
Model: "sequential_9"

Layer (type)	Output Shape	Param #
flatten_9 (Flatten)	(None, 784)	0
dense_43 (Dense)	(None, 1000)	785000
dense_44 (Dense)	(None, 1000)	1001000
dense_45 (Dense)	(None, 1000)	1001000
dense_46 (Dense)	(None, 1000)	1001000
dense_47 (Dense)	(None, 1000)	1001000
dense_48 (Dense)	(None, 10)	10010

Total params: 4,799,010
Trainable params: 4,799,010
Non-trainable params: 0

None



Train Accuracy = 0.970
Train Loss = 0.102

Validation Accuracy = 0.976
Validation Loss = 0.091

Test Accuracy = 0.967
Test Loss = 0.105

Taining Time = 7.548 second

8. Adjust the learning rate. Try a value of 0.0001. Does it make a difference?

Does it make a difference?

Hint : default learning rate is 0.01 

validation Accuracy decreased from 0.987 to 0.976, testing Accuracy also decreased from 0.979 to 0.967

Q9: (Adjust the Learning rate)



Model: (use Learning rate 0.02)

HyperParameters:
BUFFER SIZE =10000, BATCH SIZE = 100, NUM_EPOCHS = 5
Model: "sequential_10"

Layer (type)	Output Shape	Param #
flatten_10 (Flatten)	(None, 784)	0
dense_49 (Dense)	(None, 1000)	785000
dense_50 (Dense)	(None, 1000)	1001000
dense_51 (Dense)	(None, 1000)	1001000
dense_52 (Dense)	(None, 1000)	1001000
dense_53 (Dense)	(None, 1000)	1001000
dense_54 (Dense)	(None, 10)	10010

Total params: 4,799,010
Trainable params: 4,799,010
Non-trainable params: 0

None



Train Accuracy = 0.891
Train Loss = 0.429

Validation Accuracy = 0.904
Validation Loss = 0.362

Test Accuracy = 0.901
Test Loss = 0.370

Taining Time = 10.235 second

9. Adjust the learning rate. Try a value of 0.02. Does it make a difference?

Does it make a difference?

Hint : default learning rate is 0.01 

validation Accuracy decreased from 0.987 to 0.904, testing Accuracy also decreased from 0.979 to 0.901

Q10: (Combining all together)



Model: (Customized)

HyperParameters:
BUFFER SIZE =10000, BATCH SIZE = 100, NUM_EPOCHS = 50
Model: "sequential_12"

Layer (type)	Output Shape	Param #
flatten_12 (Flatten)	(None, 784)	0
dense_61 (Dense)	(None, 600)	471000
dense_62 (Dense)	(None, 600)	360600
dense_63 (Dense)	(None, 600)	360600
dense_64 (Dense)	(None, 600)	360600
dense_65 (Dense)	(None, 600)	360600
dense_66 (Dense)	(None, 10)	6010

Total params: 1,919,410
Trainable params: 1,919,410
Non-trainable params: 0

None



Train Accuracy = 0.998
Train Loss = 0.008

Validation Accuracy = 0.996
Validation Loss = 0.015

Test Accuracy = 0.982
Test Loss = 0.145

Taining Time = 78.653 second

Final Model

Validation Accuracy = 0.998



Test Accuracy = 0.982

CNN Network

CNN: Experiment 1

```
model CNN_exp_1 = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), input_shape=(28,28,1)),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
    tf.keras.layers.Conv2D(64, (3, 3)),
    tf.keras.layers.Conv2D(64, (3, 3)),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
    tf.keras.layers.Flatten(input_shape=(28,28,1)),
    tf.keras.layers.Dense(100,activation='relu',kernel_initializer='he_uniform'),
    tf.keras.layers.Dense(output_size,activation='softmax')
])
```



```
540/540 - 5s - loss: 0.0167 - accuracy: 0.9945 - val_loss: 0.0565 - val_accuracy: 0.9862
1/1 [=====] - 1s 714ms/step - loss: 0.0497 - accuracy: 0.9875
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```



CNN: Experiment 2

```
hidden_layer_size = 128

model_CNN_exp_2 = tf.keras.Sequential([
    tf.keras.layers.Conv2D(64, (3, 3), activation = 'relu', input_shape=(28,28,1)),
    tf.keras.layers.Conv2D(64, (3, 3), activation = 'relu'),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
    tf.keras.layers.Conv2D(128, (3, 3), activation = 'relu'),
    tf.keras.layers.Conv2D(128, (3, 3), activation = 'relu'),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
    tf.keras.layers.Conv2D(256, (3, 3), activation = 'relu'),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
    tf.keras.layers.Flatten(input_shape=(28,28,1)),
    tf.keras.layers.Dense(512, activation='relu', kernel_initializer='he_uniform'),
    tf.keras.layers.Dense(output_size, activation='softmax')
])
```



CNN: Experiment 3

```
model CNN_exp_3 = tf.keras.Sequential([
    tf.keras.layers.Conv2D(64, (5, 5), activation = 'relu', input_shape=(28,28,1)),
    tf.keras.layers.Conv2D(64, (5, 5), activation = 'relu'),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),

    tf.keras.layers.Conv2D(128, (3, 3), activation = 'relu'),
    tf.keras.layers.Conv2D(128, (3, 3), activation = 'relu'),
    tf.keras.layers.Conv2D(256, (3, 3), activation = 'relu'),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
    tf.keras.layers.Flatten(input_shape=(28,28,1)),
    tf.keras.layers.Dense(512, activation='relu', kernel_initializer='he_uniform'),
    tf.keras.layers.Dense(output_size, activation='softmax')
])
```

40/540 - 11s - loss: 0.0173 - accuracy: 0.9944 - val_loss: 0.0284 - val_accuracy: 0.9902
/1 [=====] - 2s 2s/step - loss: 0.0286 - accuracy: 0.9915
ict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

