
IPCV Part III miniproject 2

Last updated on 2020-05-21

SESSION 1:

- This assignment is centered in the use of [Keras](#)
- Keras is an open-source neural-network library. Designed to enable fast experimentation with deep neural networks, it focuses on being user-friendly, modular, and extensible.
- Keras itself can work using different motors. We will use it with TensorFlow under the hood.
- We will analyze the vanishing gradient problem and initialization methods together with other optimization methods.

```
In [1]: # TO-DO: Include your names and NIAs here:
student_data = [{'name': 'Mohamed Hassan', 'nia': 'NIA of 1st student'},
                 {'name': 'Sebastian Cajas', 'nia': 'NIA of 2nd student'}]
```

Import the libraries

TensorFlow officially included Keras, so if you have TensorFlow, you have keras!

```
In [21]: # Imports
import numpy as np
import matplotlib.pyplot as plt

import tensorflow as tf
from tensorflow import keras

import keras as k
from keras.models import Sequential
from keras.layers import Dense

from sklearn.preprocessing import StandardScaler
```

```
In [3]: !git clone https://github.com/luisferuam/DLFBT-LAB
import sys
sys.path.append('DLFBT-LAB')
import dlfbt
```

```
Cloning into 'DLFBT-LAB'...
remote: Enumerating objects: 99, done.
remote: Counting objects: 100% (99/99), done.
remote: Compressing objects: 100% (73/73), done.
remote: Total 99 (delta 49), reused 71 (delta 24), pack-reused 0
Unpacking objects: 100% (99/99), done.
```

Data set

```
In [22]: dataset_url = 'https://raw.githubusercontent.com/jbrownlee/Datasets/master/phone
# Details https://raw.githubusercontent.com/jbrownlee/Datasets/master/phoneme.na

dataset = np.loadtxt(dataset_url, delimiter=',')

# Split database in attributes and classes
print("dataset shape =", dataset.shape)
#-----
# T0-D0 block: Divide attributes and classes/labels. Store the number of attributes
#-----
x = dataset[:, :-1]
y = dataset[:, -1:]
x_size = x.shape
print("features shape =", x.shape)
print("labels shape =", y.shape)
#-----
# End of T0-D0 block
#-----

dataset shape = (5404, 6)
features shape = (5404, 5)
labels shape = (5404, 1)
```

```
In [5]: # Final result with the classes stored in y
print(x_size)
print(x)
print(y)
```

```
(5404, 5)
[[ 1.24  0.875 -0.205 -0.078  0.067]
 [ 0.268  1.352  1.035 -0.332  0.217]
 [ 1.567  0.867  1.3    1.041  0.559]
 ...
 [ 1.031  0.584  1.866  1.532 -0.671]
 [ 0.15   0.933  2.363 -0.742 -0.617]
 [ 0.137  0.714  1.35   0.972 -0.63 ]]
[[0.]
 [0.]
 [0.]
 ...
 [1.]
 [0.]
 [1.]]
```

```
In [6]: np.unique(y)
```

```
Out[6]: array([0., 1.])
```

```
In [23]: # Normalize the data
scaler = StandardScaler()
scaler.fit(x)
x = scaler.transform(x)
```

Defining our model

- From the input to the output in keras we can define the properties of each layer (size, activation function, connectivity topology...) with the sequential mode.
- In this case we are going to create our basic multilayer feedforward network

```
In [8]: # Define the model using keras
nn = Sequential()

#-----
# T0-D0 block: Add fully connected layers to create a MLP like in assignment 1
#-----
input_size = 5
output_size = 1
hidden_layer_size = 100

nn = tf.keras.Sequential([
    tf.keras.layers.Dense(input_size, input_dim=input_size, activation="relu"),
    tf.keras.layers.Dense(hidden_layer_size, activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size, activation='relu'),
    tf.keras.layers.Dense(output_size, activation='sigmoid')
])
#-----
# End of T0-D0 block
#-----
```

Compile the network

- Compile is the step where our network is created
- Here we have to define different aspects involved in the training of the network
- In each section you have an URL to the official documentation. Take a look at the availability of different strategies in each case.
- It is possible to also define your own functions for this.

Optimizer

Strategy to calculate the weights corrections

<https://keras.io/api/optimizers/>

Loss function

The purpose of loss functions is to compute the quantity that a model should seek to minimize during training.

<https://keras.io/api/losses/>

Metrics (results)

A metric is a function that is used to judge the performance of your model.

Metric functions are similar to loss functions, except that the results from evaluating a metric are not used when training the model. Note that you may use any loss function as a metric.

<https://keras.io/api/metrics/>

Initialization

Initializers define the way to set the initial weights of your network layers.

<https://keras.io/api/layers/initializers/>

```
In [ ]: # Compile
#-----
# T0-D0 block: Compile your network, to reproduce the assignment 1 MLP
#-----
opt = tf.keras.optimizers.SGD(learning_rate=0.001)
nn.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])
#-----
# End of T0-D0 block
#-----
```

Train the network

The Fit method trains the network according to the data.

Here we introduce all the data together and select a 20% of the data for validation purposes.

Other ways to do this are allowed, including the optimization of the parameters.

https://keras.io/api/models/model_training_apis/#fit-method

```
In [ ]: tf.keras.backend.set_floatx('float64')
```

```
In [ ]: # Fit
history = nn.fit(x, y, epochs=500, verbose=2, validation_split=0.2)

# Fit
#history = nn.fit(x, y, epochs=500, verbose=0, validation_split=0.2)
```

```
Epoch 1/500
136/136 - 4s - loss: 0.6276 - accuracy: 0.7030 - val_loss: 0.6153 - val_accuracy: 0.7206
Epoch 2/500
136/136 - 0s - loss: 0.6143 - accuracy: 0.7030 - val_loss: 0.6020 - val_accuracy: 0.7206
Epoch 3/500
136/136 - 0s - loss: 0.6032 - accuracy: 0.7030 - val_loss: 0.5912 - val_accuracy: 0.7206
Epoch 4/500
136/136 - 0s - loss: 0.5942 - accuracy: 0.7030 - val_loss: 0.5819 - val_accuracy: 0.7206
Epoch 5/500
136/136 - 0s - loss: 0.5865 - accuracy: 0.7030 - val_loss: 0.5741 - val_accuracy: 0.7206
Epoch 6/500
136/136 - 0s - loss: 0.5800 - accuracy: 0.7030 - val_loss: 0.5674 - val_accuracy: 0.7206
Epoch 7/500
136/136 - 0s - loss: 0.5745 - accuracy: 0.7030 - val_loss: 0.5617 - val_accuracy: 0.7206
Epoch 8/500
136/136 - 0s - loss: 0.5697 - accuracy: 0.7030 - val_loss: 0.5569 - val_accuracy: 0.7206
```

```

y: 0.8113
Epoch 486/500
136/136 - 0s - loss: 0.3834 - accuracy: 0.8075 - val_loss: 0.3814 - val_accurac
y: 0.8113
Epoch 487/500
136/136 - 0s - loss: 0.3832 - accuracy: 0.8075 - val_loss: 0.3813 - val_accurac
y: 0.8113
Epoch 488/500
136/136 - 0s - loss: 0.3831 - accuracy: 0.8082 - val_loss: 0.3812 - val_accurac
y: 0.8113
Epoch 489/500
136/136 - 0s - loss: 0.3829 - accuracy: 0.8073 - val_loss: 0.3811 - val_accurac
y: 0.8122
Epoch 490/500
136/136 - 0s - loss: 0.3828 - accuracy: 0.8082 - val_loss: 0.3810 - val_accurac
y: 0.8122
Epoch 491/500
136/136 - 0s - loss: 0.3826 - accuracy: 0.8082 - val_loss: 0.3809 - val_accurac
y: 0.8122
Epoch 492/500
136/136 - 0s - loss: 0.3825 - accuracy: 0.8089 - val_loss: 0.3808 - val_accurac
y: 0.8122
Epoch 493/500
136/136 - 0s - loss: 0.3824 - accuracy: 0.8094 - val_loss: 0.3808 - val_accurac
y: 0.8131
Epoch 494/500
136/136 - 0s - loss: 0.3822 - accuracy: 0.8092 - val_loss: 0.3808 - val_accurac
y: 0.8131
Epoch 495/500
136/136 - 0s - loss: 0.3821 - accuracy: 0.8087 - val_loss: 0.3806 - val_accurac
y: 0.8131
Epoch 496/500
136/136 - 0s - loss: 0.3819 - accuracy: 0.8094 - val_loss: 0.3804 - val_accurac
y: 0.8131
Epoch 497/500
136/136 - 0s - loss: 0.3818 - accuracy: 0.8089 - val_loss: 0.3804 - val_accurac
y: 0.8131
Epoch 498/500
136/136 - 0s - loss: 0.3817 - accuracy: 0.8087 - val_loss: 0.3801 - val_accurac
y: 0.8122
Epoch 499/500
136/136 - 0s - loss: 0.3815 - accuracy: 0.8087 - val_loss: 0.3799 - val_accurac
y: 0.8131
Epoch 500/500
136/136 - 0s - loss: 0.3814 - accuracy: 0.8089 - val_loss: 0.3799 - val_accurac
y: 0.8131

```

```

In [ ]: # Network details
nn.summary()
print('\n\n')

# Evaluate (similar to fit but just considering 1 epoch iteration without chang
loss, accuracy = nn.evaluate(x, y)
print('Accuracy: %.2f' % (accuracy*100))

```

```

# Also, the predict method is available to classify unlabeled data

```

```

Model: "sequential_1"

```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 5)	30
dense_1 (Dense)	(None, 100)	600

dense_2 (Dense)	(None, 100)	10100
dense_3 (Dense)	(None, 1)	101
=====		
Total params: 10,831		
Trainable params: 10,831		
Non-trainable params: 0		
=====		

169/169 [=====] - 0s 2ms/step - loss: 0.3809 - accuracy: 0.8101
Accuracy: 81.01

Plot data

- History object saves the different epoch data


```
In [ ]: # Plot history


print(history.history.keys())

plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation accuracy')
plt.title('Accuracy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='Validation loss')
plt.title('Mean Square Error')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()
```

dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])





TO-DO block: Explain what you observe

Using a small number of neurons in the hidden layer (\sim input layer)

When using a small number of neurons, smaller/equal than the input layer, the accuracy is not very high unless the number of epochs is increased. Using a 5 neurons for example and 5 epochs, the accuracy remains stable at 66%, while if the epochs is increased to 500, the accuracy increases up to 88%. From a bigger number of epochs we can also observe a smoother plot

Using a big number of neurons in the hidden layer (\gg input layer)

There will be overfitting, increasing the number of layers up to 500 and 8 layers depth, of almost a 6%.

End of TO-DO block

Optimize the network design

- Change the network architecture, introducing more layers and neurons to obtain a better result.

You can:

- Add more and different type of layers
- Change the activation functions
- Change the loss / optimizer
- Change your initialization

```
In [8]: from keras.optimizers import SGD
#-----
# TO-DO block: Include your code below
#-----
input_size = 5
output_size = 1
hidden_layer_size = 500

nn = tf.keras.Sequential([ # Sequential means a linear stack of layers
    tf.keras.layers.Dense(hidden_layer_size, input_dim=5, activation='relu', kernel_initializer='he_normal'),
    tf.keras.layers.Dense(hidden_layer_size, activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size, activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size, activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size, activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size, activation='tanh'),
    tf.keras.layers.Dense(output_size, activation='sigmoid')
])

opt = SGD(lr=0.01, momentum=0.9)

nn.compile(loss='binary_crossentropy', optimizer=opt, metrics=['accuracy'])

#-----
# End of TO-DO block
#-----

# Fit
history = nn.fit(x, y, epochs=250, verbose=0, validation_split=0.2)

# Plot history
print(history.history.keys())

plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation accuracy')
plt.title('Accuracy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()
```



```



plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='Validation loss')
plt.title('Loss using binary cross-entropy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

# Network details
nn.summary()
print('\n\n')

# Evaluate (similar to fit but just 1 epoch iteration without changing the network)
loss, accuracy = nn.evaluate(x, y)
print('Accuracy: %.2f' % (accuracy*100))

```

dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])



Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_14 (Dense)	(None, 500)	3000
dense_15 (Dense)	(None, 500)	250500

dense_16 (Dense)	(None, 500)	250500
dense_17 (Dense)	(None, 500)	250500
dense_18 (Dense)	(None, 500)	250500
dense_19 (Dense)	(None, 500)	250500
dense_20 (Dense)	(None, 1)	501
=====		
Total params: 1,256,001		
Trainable params: 1,256,001		
Non-trainable params: 0		
=====		

169/169 [=====] - 0s 3ms/step - loss: 0.1239 - accuracy: 0.9669
Accuracy: 96.69

SESSION 2

Dataset input

```
In [28]: # Load here your selected dataset considering input and output dimensions
import tensorflow_datasets as tfds
mnist_dataset,mnist_info=tfds.load(name='mnist',with_info=True, as_supervised=True)
```

```
In [29]: mnist_train,mnist_test = mnist_dataset['train'],mnist_dataset['test']
num_validation_samples = 0.1*mnist_info.splits['train'].num_examples
num_validation_samples = tf.cast(num_validation_samples,tf.int64)

num_test_samples = mnist_info.splits['test'].num_examples
num_test_samples = tf.cast(num_test_samples,tf.int64)
```

```
In [32]: def scale(image,label):
        image=tf.cast(image,tf.float32)
        image/=255.
        return image,label

scaled_train_and_validation_data=mnist_train.map(scale)
scaled_test_data = mnist_test.map(scale)

BUFFER_SIZE=10000
BATCH_SIZE=300
input_size = 784
output_size = 10
hidden_layer_size = 300
hidden_layer_size2 = 150
hidden_layer_size2 = 150
hidden_layer_size2 = 150

shuffled_train_and_validation_data = scaled_train_and_validation_data.shuffle(BUFFER_SIZE)

validation_data = shuffled_train_and_validation_data.take(num_validation_samples)
train_data = shuffled_train_and_validation_data.skip(num_validation_samples)

train_data=train_data.batch(BATCH_SIZE)
```

```
validation_data = validation_data.batch(num_validation_samples)
test_data = scaled_test_data.batch(num_test_samples)

validation_inputs, validation_targets = next(iter(validation_data))
```

Standar choice

- Evaluate the training with a standard choice of cost and activation functions, learning rate, weight initialization and network topology.
- Generate the loss and accuracy figures

```
In [33]: model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28,1)),
    tf.keras.layers.Dense(hidden_layer_size,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(output_size,activation='softmax')
])



opt = keras.optimizers.Nadam(learning_rate=0.001)
model.compile(optimizer='Adam',loss='sparse_categorical_crossentropy',metrics=['

NUM_EPOCHS= 10
history = model.fit(train_data,epochs=NUM_EPOCHS,validation_data=(validation_inp
```

```
In [34]: plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation accuracy')
plt.title('Accuracy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='Validation loss')
plt.title('Mean Square Error')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

# Evaluate (similar to fit but just 1 epoch iteration without changing the network)
loss, accuracy = model.evaluate(test_data)
print('Accuracy: %.2f' % (accuracy*100))
```



1/1 [=====] - 1s 1s/step - loss: 0.0765 - accuracy: 0.9
807
Accuracy: 98.07

Optimizations

- Evaluate each of the following optimization methods using the same representation and the duration of the training in terms of epochs to reach a choice of error and also in terms of time taken.
 - a) Regularization
 - b) Dropout
 - c) Stochastic gradient descent
 - d) Momentum (including Nesterov version)
 - e) AdaGrad
 - f) RMSProp

g) Adam

h) Optimize

```
In [ ]: # a) Regularization
# Code
model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28,1)),
    tf.keras.layers.Dense(hidden_layer_size,activation='relu',kernel_regularizer=
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu',kernel_regularizer=
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu',kernel_regularizer=
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu',kernel_regularizer=
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu',kernel_regularizer=
    tf.keras.layers.Dense(output_size,activation='softmax')
])


opt = keras.optimizers.Nadam(learning_rate=0.001)
model.compile(optimizer='Adam',loss='sparse_categorical_crossentropy',metrics=['

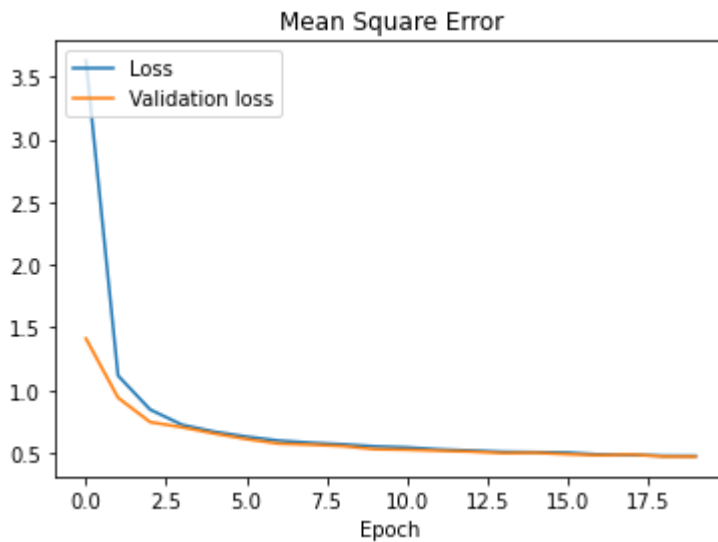
NUM_EPOCHS= 20
history = model.fit(train_data,epochs=NUM_EPOCHS,validation_data=(validation_inp
```

```
In [ ]: plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation accuracy')
plt.title('Accuracy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='Validation loss')
plt.title('Mean Square Error')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

# Evaluate (similar to fit but just 1 epoch iteration without changing the network
loss, accuracy = model.evaluate(test_data)
print('Accuracy: %.2f' % (accuracy*100))
```





1/1 [=====] - 0s 433ms/step - loss: 0.4708 - accuracy: 0.9587
Accuracy: 95.87

```
In [ ]: # b) Dropout
# Code
model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28,1)),
    tf.keras.layers.Dense(hidden_layer_size,activation='relu'),
    tf.keras.layers.Dropout(0.25),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dropout(0.25),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dropout(0.25),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dropout(0.25),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(output_size,activation='softmax')
])

opt = keras.optimizers.Nadam(learning_rate=0.001)
model.compile(optimizer='Adam',loss='sparse_categorical_crossentropy',metrics=['



NUM_EPOCHS= 20
history = model.fit(train_data,epochs=NUM_EPOCHS,validation_data=(validation_inp
```

```
In [ ]: plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation accuracy')
plt.title('Accuracy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='Validation loss')
plt.title('Mean Square Error')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

# Evaluate (similar to fit but just 1 epoch iteration without changing the network)
```

```
loss, accuracy = model.evaluate(test_data)
print('Accuracy: %.2f' % (accuracy*100))
```



1/1 [=====] - 0s 446ms/step - loss: 0.0735 - accuracy: 0.9808
Accuracy: 98.08

```
In [ ]: # c) Stochastic gradient descent
# Code
model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28,1)),
    tf.keras.layers.Dense(hidden_layer_size,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(output_size,activation='softmax')
])

opt = keras.optimizers.SGD(learning_rate=0.001)
model.compile(optimizer=opt,loss='sparse_categorical_crossentropy',metrics=['acc'])

NUM_EPOCHS= 20
history = model.fit(train_data,epochs=NUM_EPOCHS,validation_data=(validation_inp
```

```
In [ ]: plt.plot(history.history['accuracy'], label='Accuracy')
```



```

plt.plot(history.history['val_accuracy'], label='Validation accuracy')
plt.title('Accuracy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='Validation loss')
plt.title('Mean Square Error')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

# Evaluate (similar to fit but just 1 epoch iteration without changing the network)
loss, accuracy = model.evaluate(test_data)
print('Accuracy: %.2f' % (accuracy*100))

```



```

1/1 [=====] - 0s 451ms/step - loss: 0.8719 - accuracy: 0.8014
Accuracy: 80.14

```

```

In [ ]: # d) Momentum (including Nesterov version)
# Code
model = tf.keras.Sequential([

```



```

tf.keras.layers.Flatten(input_shape=(28,28,1)),
tf.keras.layers.Dense(hidden_layer_size,activation='relu'),
tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
tf.keras.layers.Dense(output_size,activation='softmax')
])

opt = keras.optimizers.SGD(learning_rate=0.001, momentum=0.9)
model.compile(optimizer= opt,loss='sparse_categorical_crossentropy',metrics=['ac

NUM_EPOCHS= 20
history = model.fit(train_data,epochs=NUM_EPOCHS,validation_data=(validation_inp

```


```


In [ ]: plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation accuracy')
plt.title('Accuracy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='Validation loss')
plt.title('Mean Square Error')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

# Evaluate (similar to fit but just 1 epoch iteration without changing the network)
loss, accuracy = model.evaluate(test_data)
print('Accuracy: %.2f' % (accuracy*100))

```





1/1 [=====] - 0s 442ms/step - loss: 0.1690 - accuracy: 0.9511
Accuracy: 95.11

```
In [ ]: # e) AdaGrad
# Code
model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28,1)),
    tf.keras.layers.Dense(hidden_layer_size,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(output_size,activation='softmax')
])



opt = keras.optimizers.Adagrad(learning_rate=0.001)
model.compile(optimizer=opt,loss='sparse_categorical_crossentropy',metrics=['acc

NUM_EPOCHS= 20
history = model.fit(train_data,epochs=NUM_EPOCHS,validation_data=(validation_inp
```

```
In [ ]: plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation accuracy')
plt.title('Accuracy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='Validation loss')
plt.title('Mean Square Error')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

# Evaluate (similar to fit but just 1 epoch iteration without changing the network)
loss, accuracy = model.evaluate(test_data)
print('Accuracy: %.2f' % (accuracy*100))
```



1/1 [=====] - 0s 426ms/step - loss: 0.3358 - accuracy: 0.9017
Accuracy: 90.17

```
In [ ]: # f) RMSProp
# Code

model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28,1)),
    tf.keras.layers.Dense(hidden_layer_size,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(output_size,activation='softmax')
])

opt = keras.optimizers.RMSprop(learning_rate=0.001)
model.compile(optimizer=opt,loss='sparse_categorical_crossentropy',metrics=['acc



NUM_EPOCHS= 20
history = model.fit(train_data,epochs=NUM_EPOCHS,validation_data=(validation_inp

In [ ]: plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation accuracy')
```

```
plt.title('Accuracy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='Validation loss')
plt.title('Mean Square Error')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

# Evaluate (similar to fit but just 1 epoch iteration without changing the network)
loss, accuracy = model.evaluate(test_data)
print('Accuracy: %.2f' % (accuracy*100))
```



```
1/1 [=====] - 0s 455ms/step - loss: 0.1286 - accuracy: 0.9809
Accuracy: 98.09
```

```
In [ ]: # g) Adam
# Code
model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28,1)),
```

```

tf.keras.layers.Dense(hidden_layer_size,activation='relu'),
tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
tf.keras.layers.Dense(output_size,activation='softmax')
])

opt = keras.optimizers.Adam(learning_rate=0.001)
model.compile(optimizer=opt,loss='sparse_categorical_crossentropy',metrics=['acc

NUM_EPOCHS= 20
history = model.fit(train_data,epochs=NUM_EPOCHS,validation_data=(validation_inp

```


```


In [ ]: plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation accuracy')
plt.title('Accuracy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='Validation loss')
plt.title('Mean Square Error')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

# Evaluate (similar to fit but just 1 epoch iteration without changing the network
loss, accuracy = model.evaluate(test_data)
print('Accuracy: %.2f' % (accuracy*100))

```





1/1 [=====] - 0s 433ms/step - loss: 0.0895 - accuracy: 0.9799
Accuracy: 97.99

```
In [ ]: # h) Optimize
# Code
model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28,28,1)),
    tf.keras.layers.Dense(hidden_layer_size,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
    tf.keras.layers.Dense(output_size,activation='softmax')
])

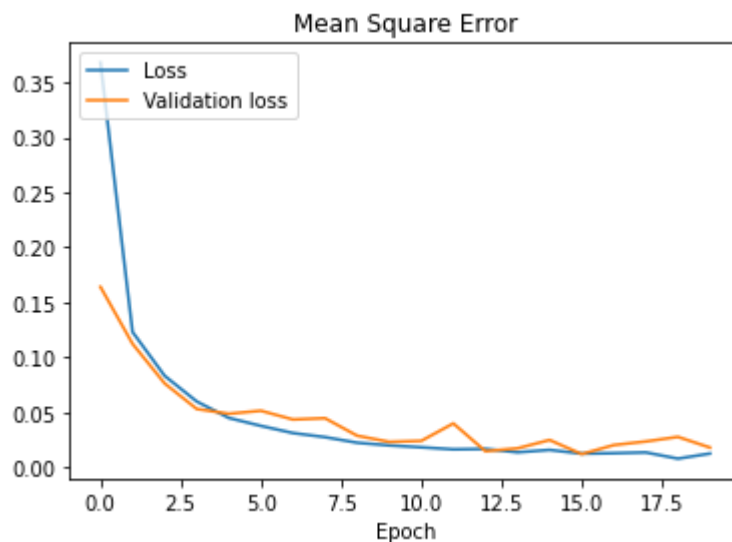

opt = keras.optimizers.Nadam(learning_rate=0.001)
model.compile(optimizer='Adam',loss='sparse_categorical_crossentropy',metrics=['

NUM_EPOCHS= 20
history = model.fit(train_data,epochs=NUM_EPOCHS,validation_data=(validation_inp
```

```
In [ ]: plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation accuracy')
plt.title('Accuracy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='Validation loss')
plt.title('Mean Square Error')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

# Evaluate (similar to fit but just 1 epoch iteration without changing the network)
loss, accuracy = model.evaluate(test_data)
print('Accuracy: %.2f' % (accuracy*100))
```



1/1 [=====] - 0s 440ms/step - loss: 0.1051 - accuracy: 0.9797
 Accuracy: 97.97

Implement parallelization in Keras with the best optimization

- Run tests for GPU presence
- Use your best implementation from the previous exercise here and compare the training time both with GPUs and without GPUs.

Comment all your results.

```
In [23]: %tensorflow_version 2.x
import tensorflow as tf
print("Tensorflow version " + tf.__version__)

try:
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver() # TPU detection
    print('Running on TPU ', tpu.cluster_spec().as_dict()['worker'])
except ValueError:
    raise BaseException('ERROR: Not connected to a TPU runtime; please see the pre

tf.config.experimental_connect_to_cluster(tpu)
```

```
tf.tpu.experimental.initialize_tpu_system(tpu)
tpu_strategy = tf.distribute.experimental.TPUStrategy(tpu)
```

Tensorflow version 2.4.1

Running on TPU ['10.16.57.162:8470']

WARNING:tensorflow:TPU system grpc://10.16.57.162:8470 has already been initialized. Reinitializing the TPU can cause previously created variables on TPU to be lost.

WARNING:tensorflow:TPU system grpc://10.16.57.162:8470 has already been initialized. Reinitializing the TPU can cause previously created variables on TPU to be lost.

INFO:tensorflow:Initializing the TPU system: grpc://10.16.57.162:8470

INFO:tensorflow:Initializing the TPU system: grpc://10.16.57.162:8470

INFO:tensorflow:Clearing out eager caches

INFO:tensorflow:Clearing out eager caches

INFO:tensorflow:Finished initializing TPU system.

INFO:tensorflow:Finished initializing TPU system.

WARNING:absl: `tf.distribute.experimental.TPUStrategy` is deprecated, please use the non experimental symbol `tf.distribute.TPUStrategy` instead.

INFO:tensorflow:Found TPU system:

INFO:tensorflow:Found TPU system:

INFO:tensorflow:*** Num TPU Cores: 8

INFO:tensorflow:*** Num TPU Cores: 8

INFO:tensorflow:*** Num TPU Workers: 1

INFO:tensorflow:*** Num TPU Workers: 1

INFO:tensorflow:*** Num TPU Cores Per Worker: 8

INFO:tensorflow:*** Num TPU Cores Per Worker: 8

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:localhost/replica:0/task:0/device:CPU:0, CPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:localhost/replica:0/task:0/device:CPU:0, CPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:CPU:0, CPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:CPU:0, CPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:0, TPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:0, TPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:1, TPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:1, TPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:2, TPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:2, TPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:3, TPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:3, TPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:4, TPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:4, TPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:5, TPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:5, TPU, 0, 0)

INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:5, TPU, 0, 0)


```
sk:0/device:TPU:6, TPU, 0, 0)
```

```
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:6, TPU, 0, 0)
```

```
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:7, TPU, 0, 0)
```

```
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:7, TPU, 0, 0)
```

```
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU_SYSTEM:0, TPU_SYSTEM, 0, 0)
```

```
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU_SYSTEM:0, TPU_SYSTEM, 0, 0)
```

```
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:XLA_CPU:0, XLA_CPU, 0, 0)
```

```
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:XLA_CPU:0, XLA_CPU, 0, 0)
```

Guess the year comparason with CPU - GPU - TPU

```
In [26]: input_size = 5
output_size = 1
hidden_layer_size = 500

nn = tf.keras.Sequential([ # Sequential means a linear stack of layers
    tf.keras.layers.Dense(hidden_layer_size,input_dim=5,activation='relu', kernel_initializer='glorot_uniform'),
    tf.keras.layers.Dense(hidden_layer_size,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size,activation='relu'),
    tf.keras.layers.Dense(hidden_layer_size,activation='tanh'),
    tf.keras.layers.Dense(output_size,activation='sigmoid')
])

opt = SGD(lr=0.01, momentum=0.9)

nn.compile(loss='binary_crossentropy', optimizer=opt,metrics=['accuracy'])

#-----
# End of TO-DO block
#-----
import time

s = time.time()
# Fit
history = nn.fit(x, y, epochs=5,batch_size=1024, verbose=1, validation_split=0.2)
e = time.time()

print(e - s)
# Plot history

print(history.history.keys())

plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation accuracy')
plt.title('Accuracy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()
```

```

plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='Validation loss')
plt.title('Loss using binary cross-entropy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

# Network details
nn.summary()
print('\n\n')


# Evaluate (similar to fit but just 1 epoch iteration without changing the network)
loss, accuracy = nn.evaluate(x, y)
print('Accuracy: %.2f' % (accuracy*100))


```

```

Epoch 1/5
5/5 [=====] - 1s 180ms/step - loss: 0.6256 - accuracy:
0.6773 - val_loss: 0.5084 - val_accuracy: 0.7206
Epoch 2/5
5/5 [=====] - 1s 132ms/step - loss: 0.5008 - accuracy:
0.7288 - val_loss: 0.4527 - val_accuracy: 0.7956
Epoch 3/5
5/5 [=====] - 1s 137ms/step - loss: 0.4515 - accuracy:
0.7943 - val_loss: 0.4290 - val_accuracy: 0.7882
Epoch 4/5
5/5 [=====] - 1s 134ms/step - loss: 0.4370 - accuracy:
0.7824 - val_loss: 0.4119 - val_accuracy: 0.8057
Epoch 5/5
5/5 [=====] - 1s 132ms/step - loss: 0.4033 - accuracy:
0.8123 - val_loss: 0.3996 - val_accuracy: 0.8020
4.068118333816528
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

```





Model: "sequential_13"

Layer (type)	Output Shape	Param #
dense_76 (Dense)	(None, 500)	3000
dense_77 (Dense)	(None, 500)	250500
dense_78 (Dense)	(None, 500)	250500
dense_79 (Dense)	(None, 500)	250500
dense_80 (Dense)	(None, 500)	250500
dense_81 (Dense)	(None, 500)	250500
dense_82 (Dense)	(None, 1)	501
Total params: 1,256,001		
Trainable params: 1,256,001		
Non-trainable params: 0		

169/169 [=====] - 1s 4ms/step - loss: 0.3992 - accuracy: 0.8077
Accuracy: 80.77

Computational time between CPU - GPU - TPU
for best model - guess the year dataset

Epoch time

Batch Size	CPU	GPU	TPU
256	5.038s	1.60s	1.63s
512	4.430s	1.35s	1.98s
1024	4.06s	1.258s	1.635s

Step time per epoch

Batch Size	CPU	GPU	TPU
256	1s/51 μ s	0s/14 μ s	1s/22 μ s
512	1s/80 μ s	0s/14 μ s	1s/24 μ s
1024	1s/133 μ s	0s/19 μ s	0/42 μ s

Mnist comparason with CPU - GPU

```
In [35]: from keras.optimizers import SGD
import timeit
import time

#from tensorflow.contrib.tpu.python.tpu import keras_support

def get_model():

    model = tf.keras.Sequential([
        tf.keras.layers.Flatten(input_shape=(28,28,1)),
        tf.keras.layers.Dense(hidden_layer_size,activation='relu'),
        tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
        tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
        tf.keras.layers.Dense(hidden_layer_size2,activation='relu'),
        tf.keras.layers.Dense(output_size,activation='softmax')
    ])

    opt = keras.optimizers.Nadam(learning_rate=0.001)
    model.compile(optimizer='Adam',loss='sparse_categorical_crossentropy',metrics=

    return model
# End of T0-D0 block
#-----
model = get_model()

NUM_EPOCHS= 20
s = time.time()
history = model.fit(train_data,epochs=NUM_EPOCHS,
                    batch_size=256,
                    validation_data=(validation_inputs, validation_targets),verbose=0)

#history = model.fit(train_data,epochs=NUM_EPOCHS,validation_data=(validation_in
e = time.time()

print("The computation time was: ", e - s)
# Plot history

print(history.history.keys())

plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation accuracy')
plt.title('Accuracy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
```



```
plt.show()

plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='Validation loss')
plt.title('Loss using binary cross-entropy')
plt.ylabel('')
plt.xlabel('Epoch')
plt.legend(loc="upper left")
plt.show()

# Network details
model.summary()
print('\n\n')

# Evaluate (similar to fit but just 1 epoch iteration without changing the network)
loss, accuracy = model.evaluate(test_data)
print('Accuracy: %.2f' % (accuracy*100))
```

The computation time was: 98.12066292762756
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])



Model: "sequential_18"

Layer (type)	Output Shape	Param #
flatten_15 (Flatten)	(None, 784)	0

dense_107 (Dense)	(None, 300)	235500
dense_108 (Dense)	(None, 150)	45150
dense_109 (Dense)	(None, 150)	22650
dense_110 (Dense)	(None, 150)	22650
dense_111 (Dense)	(None, 150)	22650
dense_112 (Dense)	(None, 10)	1510
=====		
Total params: 350,110		
Trainable params: 350,110		
Non-trainable params: 0		

1/1 [=====] - 1s 655ms/step - loss: 0.0894 - accuracy: 0.9798
Accuracy: 97.98

Computational time between CPU - GPU - TPU for best model - MNIST model

Epoch time

Batch Size	CPU	GPU
256	104.6067s	75.24s
512	102s	74.6s
1024	103.75s	80.05s

Step time per epoch

Batch Size	CPU	GPU
256	5.2s/25.5μs	4s/17μs
512	5.1s/24μs	4s/17μs
1024	5s/24μs	4s/17μs

Convolutional Neural Networks

Implement a convolutional neural network and apply it to classify the images of the CIFAR-10 dataset. The network should have at least the following characteristics:

- Convolutional layers.
- Pooling layers.
- Some regularization mechanism, such as dropout or L2 regularization.

Batch normalization layers.

The network should obtain at least 75% accuracy on the test set.

```
In [ ]: ##tensorflow_version 2.x
import tensorflow as tf
from tensorflow import keras
from keras import backend as K
from keras.utils.vis_utils import plot_model

import numpy as np
import matplotlib.pyplot as plt
```

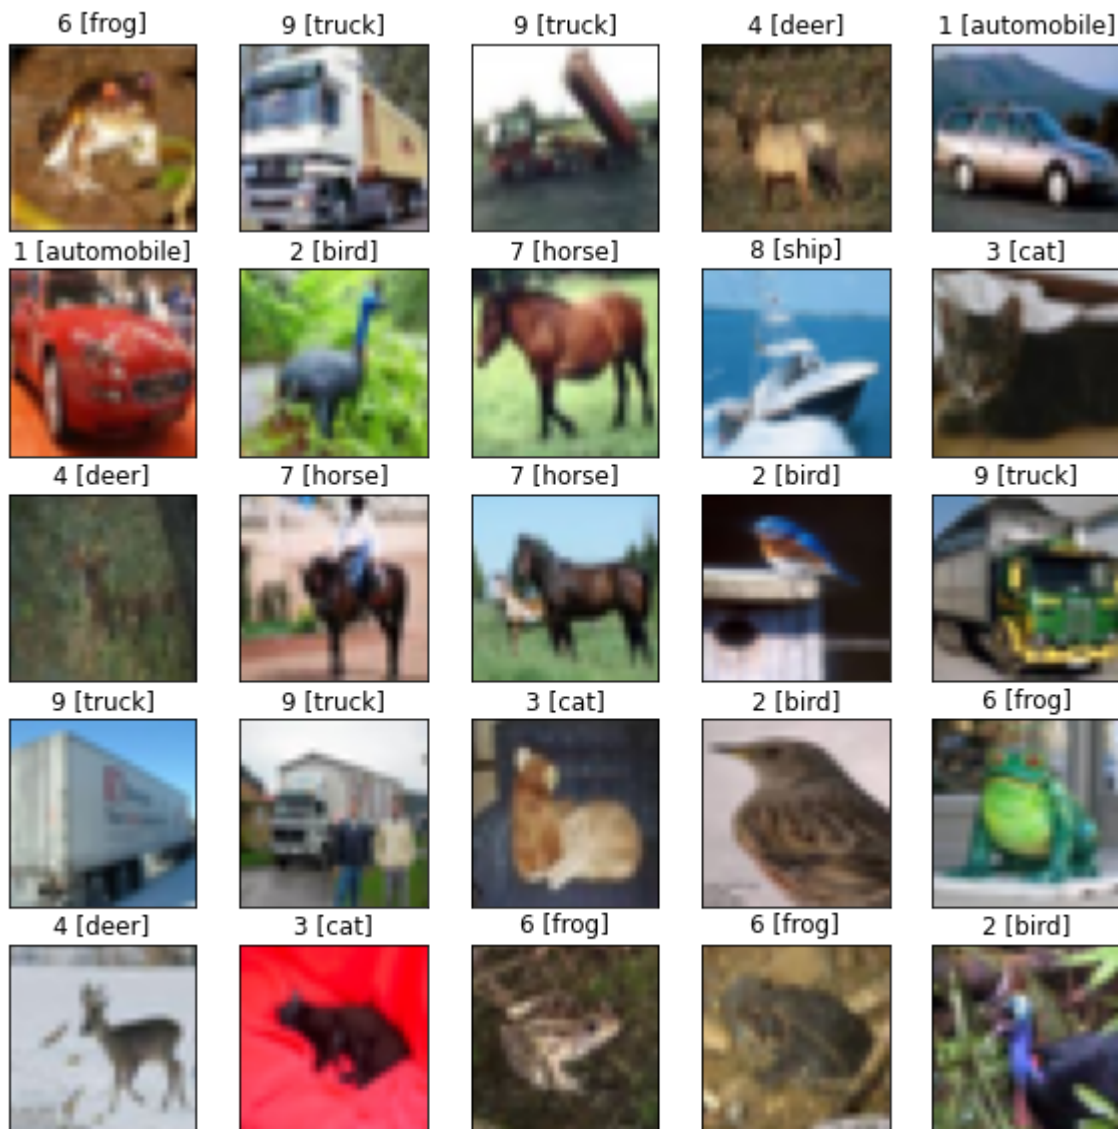
```
In [ ]: (train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.cifar100
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

print(train_images.shape)
print(train_labels.shape)
for t in train_labels[:10]:
    print(t[0], class_names[t[0]])

print(test_images.shape)
print(test_labels.shape)
for t in test_labels[:10]:
    print(t[0], class_names[t[0]])
```

```
(50000, 32, 32, 3)
(50000, 1)
6 frog
9 truck
9 truck
4 deer
1 automobile
1 automobile
2 bird
7 horse
8 ship
3 cat
(10000, 32, 32, 3)
(10000, 1)
3 cat
8 ship
8 ship
0 airplane
6 frog
6 frog
1 automobile
6 frog
3 cat
1 automobile
```

```
In [ ]: plt.figure(figsize=(10, 10))
for i in range(25):
    plt.subplot(5, 5, i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(train_images[i])
    plt.title("%d [%s]" % (train_labels[i][0], class_names[train_labels[i][0]]))
```



```
In [ ]: train_images.mean(axis=0).shape
```

```
Out[ ]: (32, 32, 3)
```

Normalization: Data - column-wise mean

```
In [ ]: train_images_orig = train_images
test_images_orig = test_images
mean_img = train_images.mean(axis=0) # axis 0 (columns), axis 1 (rows) -> shape
train_images = train_images - mean_img
test_images = test_images - mean_img
```

Hot encoding labels

```
In [ ]: train_labels_one_hot = keras.utils.to_categorical(train_labels, 10)
test_labels_one_hot = keras.utils.to_categorical(test_labels, 10)
train_labels_one_hot[1:5]
```

```
Out[ ]: array([[0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],
               [0., 0., 0., 0., 0., 0., 0., 0., 0., 1.],
               [0., 0., 0., 0., 1., 0., 0., 0., 0., 0.],
               [0., 1., 0., 0., 0., 0., 0., 0., 0., 0.]], dtype=float32)
```


Fully connected network

```
In [ ]: from keras.preprocessing.image import ImageDataGenerator

datagen = ImageDataGenerator(width_shift_range=0.1,
                             height_shift_range=0.1,
                             rotation_range=15,
                             horizontal_flip=True)

iter = datagen.flow(train_images_orig[:10], train_labels[:10], batch_size=10)

In [ ]: K.clear_session()

model = tf.keras.Sequential()

model.add(tf.keras.layers.Input(shape=(32, 32, 3)))

model.add(tf.keras.layers.Conv2D(32, kernel_size=3, activation='relu', padding=''))
model.add(tf.keras.layers.BatchNormalization())
model.add(tf.keras.layers.Conv2D(32, kernel_size=3, activation='relu', padding=''))
model.add(tf.keras.layers.BatchNormalization())
model.add(tf.keras.layers.MaxPool2D(pool_size=2))

model.add(tf.keras.layers.Conv2D(64, kernel_size=3, activation='relu', padding=''))
model.add(tf.keras.layers.BatchNormalization())
model.add(tf.keras.layers.Conv2D(64, kernel_size=3, activation='relu', padding=''))
model.add(tf.keras.layers.BatchNormalization())
model.add(tf.keras.layers.MaxPool2D(pool_size=2))

model.add(tf.keras.layers.Conv2D(128, kernel_size=3, activation='relu', padding=''))
model.add(tf.keras.layers.BatchNormalization())
model.add(tf.keras.layers.Conv2D(128, kernel_size=3, activation='relu', padding=''))
model.add(tf.keras.layers.BatchNormalization())
model.add(tf.keras.layers.MaxPool2D(pool_size=2))

model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(512, activation='relu'))
model.add(tf.keras.layers.BatchNormalization())
model.add(tf.keras.layers.Dropout(0.5))
model.add(tf.keras.layers.Dense(10, activation='softmax'))

print(model.summary())
plot_model(model, show_shapes=True, show_layer_names=True)

model.compile(optimizer=keras.optimizers.Adam(learning_rate=1.e-3), loss='categorical_crossentropy')
iter = datagen.flow(train_images, train_labels_one_hot, batch_size=1000)

nepochs = 50
history = model.fit(iter,
                    epochs=nepochs,
                    steps_per_epoch=50,
                    validation_data=(test_images, test_labels_one_hot))
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
batch_normalization (Batch Normalization)	(None, 32, 32, 32)	128

conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 32)	128
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
batch_normalization_2 (Batch Normalization)	(None, 16, 16, 64)	256
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
batch_normalization_3 (Batch Normalization)	(None, 16, 16, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
batch_normalization_4 (Batch Normalization)	(None, 8, 8, 128)	512
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
batch_normalization_5 (Batch Normalization)	(None, 8, 8, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 512)	1049088
batch_normalization_6 (Batch Normalization)	(None, 512)	2048
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 10)	5130

```

=====
Total params: 1,345,066
Trainable params: 1,343,146
Non-trainable params: 1,920

```

None

Epoch 1/50

50/50 [=====] - 53s 430ms/step - loss: 2.1858 - acc: 0.3301 - val_loss: 2.3637 - val_acc: 0.2366

Epoch 2/50

50/50 [=====] - 22s 431ms/step - loss: 1.4189 - acc: 0.5141 - val_loss: 1.6698 - val_acc: 0.4085

Epoch 3/50

50/50 [=====] - 22s 432ms/step - loss: 1.2260 - acc: 0.5768 - val_loss: 1.3654 - val_acc: 0.5139

Epoch 4/50

50/50 [=====] - 21s 427ms/step - loss: 1.0888 - acc: 0.6263 - val_loss: 1.1979 - val_acc: 0.5659

Epoch 5/50

50/50 [=====] - 22s 430ms/step - loss: 0.9782 - acc: 0.6638 - val_loss: 0.9816 - val_acc: 0.6538

Epoch 6/50

50/50 [=====] - 21s 428ms/step - loss: 0.9092 - acc: 0.6896 - val_loss: 1.0435 - val_acc: 0.6626

Epoch 7/50

50/50 [=====] - 22s 432ms/step - loss: 0.8468 - acc: 0.7110 - val_loss: 1.0280 - val_acc: 0.6667

Epoch 8/50

50/50 [=====] - 22s 432ms/step - loss: 0.7950 - acc: 0.7291 - val_loss: 0.7628 - val_acc: 0.7412

